


Where the interpretable object sits relative to the executed decision path

transparent layers


interpretability sits around the core



executed path: $x \rightarrow h \rightarrow a$
constraints govern a ; the mechanism
from h to a stays unnamed

post-hoc explanation


interpretability fitted outside the policy



executed path unchanged; q_ϕ read off h
after training — describes, does not
guarantee $\text{do}(b)$ semantics

born legible

interpretability inside the executed path



executed path: $x \rightarrow b \rightarrow \pi_{\text{leg}} \rightarrow a$
the object interpreted is the object executed