

CrystalRL — Crystal Clear Reinforcement Learning: from a Constrained Hierarchy to an Agent Legible by Construction

Ivan Pavliuk

Draft v0.3 — July 2026

Abstract

Post-hoc explanation of a trading policy answers the wrong question: it tells a story *about* a black box while the box goes on being one. We report a three-generation worked example of the alternative, with the measurement discipline that keeps “legible” from being a vibe. The paper traces the path as it was walked. **CHRL**, a constrained hierarchical PPO trader, factors decisions into three readable layers under five guardrails and captures 93% of benchmark Sharpe at 52% less drawdown — but its core is a 64-dimensional unnamed latent with ~ 214 tuning knobs. A dedicated post-hoc program (**Interpretable-CHRL**) discovers a complete eight-primitive behavioral vocabulary over the frozen policy and shows that targeted edits of its hidden states and actions survive matched-random controls (+0.19 to +0.30 pp control-adjusted on a frozen 2022–23 rollout) — and simultaneously measures where post-hoc explanation stops: a nine-rule story reproduces only 0.244 of the deployed policy’s behavior, no bound exists on what a memory edit can touch, and steering is a correlation, not an intervention. **CRYSTAL-1** rebuilds the core so the stopping point disappears: memory is a *named* regime posterior $b_t = \mathbb{P}(\text{bear} \mid \mathcal{F}_t)$ computed by an explicit Bayes filter, the policy is a readable rule at full performance parity, and the command surface is five named levers. On the identical scalar, $\text{CRYSTALSCORE} = \text{Faithfulness} \times \text{Simulatability} \times \text{Stability}$ at a fixed $K \leq 9$ story budget, the journey moves the agent from 0.15 to 0.92 on real market data, and the entire gap is Simulatability. The certified readable rule — “ $b_t > 0.66 \Rightarrow \text{exposure } 0.74$ ” — held one-shot on 629 untouched trading days (Sharpe 1.46 vs 1.39 buy-and-hold, max drawdown -12.9% vs -15.5%). We close with what the journey establishes about interpretable reinforcement learning in general: which axis born-legible actually buys, why faithfulness probes can silently lie on threshold policies, and where legibility is provably not free.

1 Introduction: the question behind “explainable”

When a reinforcement-learning agent trades real money, “why did it do that?” is not a curiosity — it is a regulatory demand, a debugging requirement, and (in our companion work on self-improving systems) a safety precondition: *an agent that changes itself may only be allowed to do so to the degree it can be read*. The dominant practice answers with post-hoc explanation: train a black box, then probe it. Rudin [2019] names the flaw: for high-stakes decisions, do not explain the black box; build the model interpretable in the first place. The critique is well known. What is scarcer is a *worked example with a measurement discipline*: an agent built legible by construction, evaluated on real market data against a strong black-box predecessor *and* against that predecessor’s own best post-hoc program, with the legibility difference expressed as a number all three can be scored on — and with the costs stated.

This paper is that worked example, and it has the shape of a journey because we walked it in order.¹ We first built the transparent-*layers* agent (CHRL, §2); then ran the strongest post-hoc

¹The paper’s structure deliberately follows three models: the system-paper arc of Mnih et al. [2015] (problem $\rightarrow$ architecture $\rightarrow$ measured evaluation $\rightarrow$ scope), the build-interpretable thesis of Rudin [2019], and the evaluation-rigor program of Doshi-Velez and Kim [2017], whose taxonomy CrystalScore operationalizes.

interpretability program we could over its frozen policy and measured both what it buys and where it stops (§3); then rebuilt the core so the stopping point disappears (CRYSTAL-1, §4); then scored both architectures with one scalar that could have come out against us (§6–§7). Everything that counts as evidence is computed on real data; the one synthetic environment we built is used for instrument calibration only, behind an explicit gate (§8).

Contributions:

1. **CHRL** (§2): a constrained hierarchical PPO trader whose ablation shows that *constraints, not hierarchy alone*, buy the risk-adjusted result — and a precise measurement of where its transparency stops.
2. **The post-hoc program** (§3): an eight-primitive behavioral vocabulary over the frozen policy, a taxonomy of controlled latent interventions, and control-adjusted evidence that targeted edits are real — the strongest case we can make *for* post-hoc methods, stated against our own thesis.
3. **CRYSTAL-1** (§4): named belief memory (an explicit Bayes filter), a readable-rule policy at parity, five named command levers.
4. **CrystalScore and its measurement protocol** (§6): faithfulness as intervention-following, simulatability as a parity-constrained short story, stability as seed replication — including a measurement correction (a faithfulness probe that silently lies on threshold policies) that is itself a finding.
5. **The comparison** (§7): the identical scalar on both architectures, real data only, with the behavioral-influence toolbox (reward, teachers, interventions, architecture) exercised under controls.
6. **The designed-market experiment** (§8): can a market be built that rewards only a clean mind? — with the value-of-information gate that keeps its results out of market claims.
7. **Conclusions for interpretable RL** (§11): five transferable findings and a falsifiable research program.

2 CHRL: how far do transparent layers go?

2.1 Why a constrained hierarchy

CHRL (Constrained Hierarchical RL) was our best answer to interpretability *within* the standard deep-RL recipe, and its design was an empirical response to three findings of the predecessor feature-ablation study on the same market data: (i) feature quantity is not the bottleneck — a compact macro context outperforms broad feature stacks; (ii) cash must not be treated as just another tradeable asset — doing so lets the policy hide cash-heavy behavior inside an opaque allocation vector; and (iii) the practical route to reading a trading policy is *logged, layered decisions*, not a monolithic end-to-end action. The design philosophy was “transparent behavior, not a single action vector”: factor the decision so that each stage emits an observable, auditable signal.

Concretely, decisions are factored into three layers on different rhythms, each readable on its own timescale:

layer	decision	rhythm
risk	cash vs. risky exposure	~20 trading days
group	route capital through stock clusters	inherited context
stock	buy/sell candidate selection	~5 trading days




Figure 1: CHRL. Left: the three-layer decision factorization and its five guardrails — each stage emits a loggable, auditable signal on its own rhythm. Right: four-fold walk-forward validation against the 100%-risky equal-weight benchmark: 58.8% of the return, 93.2% of the Sharpe, 52% less drawdown.

wrapped in five explicit guardrails: *exposure pacing* (the risk layer cannot flip the book daily), *confidence-scaled sizing* (weak signals trade less), *event triggers* (market stress forces earlier de-risking; recovery evidence re-opens appetite), *top-K routing* (capital flows only through the strongest candidates), and a *buy gate* (new risky exposure is blocked until recovery evidence is visible). Hierarchy as a partial legibility device has a long RL lineage [Sutton et al., 1999, Bacon et al., 2017]; the guardrails are constraints in the spirit of safe policy optimization [Achiam et al., 2017], imposed structurally rather than through the reward. The core is a PPO policy [Schulman et al., 2017] trained with the clipped surrogate

$$L^{\text{CLIP}}(\theta) = \mathbb{E}_t \left[\min(\rho_t(\theta) \hat{A}_t, \text{clip}(\rho_t(\theta), 1-\epsilon, 1+\epsilon) \hat{A}_t) \right], \quad \rho_t(\theta) = \frac{\pi_\theta(a_t | s_t)}{\pi_{\theta_{\text{old}}}(a_t | s_t)}, \quad (1)$$

whose state s_t passes through a 64-dimensional *unnamed* latent.

The ablation that justifies the name is the honest part: in the four-way comparison (flat baseline / hierarchy alone / constraints alone / full CHRL), hierarchy *alone* reduced raw performance; constraints *alone* sharply reduced drawdown and turnover; only the full combination produced the best return-to-drawdown ratio. The load-bearing ingredient is the constraint set — the hierarchy is the scaffolding that makes the constraints expressible and the decisions loggable.

2.2 What it delivers

In four-fold walk-forward validation on the real US panel, CHRL delivers mean validation return 10.05%, Sharpe 1.09, max drawdown -7.34% , mean cash weight 41.24%, and 0.84%/day L1 turnover. Against the 100%-risky equal-weight universe benchmark (return 17.09%, Sharpe 1.17, max drawdown -15.27%), it captures 93.2% of benchmark Sharpe while cutting drawdown by 52%; its return-to-drawdown ratio is 1.37 vs the benchmark’s 1.12 (Figure 1). In plain language: a deliberately constrained agent gives up return to buy a smoother ride and an auditable decision trail — exactly the trade a risk-managed mandate wants.

So the layers are transparent, the guardrails are readable, and the trader is competent. The question this paper turns on is what happens when you ask the next question: *what is the policy thinking between the layers?*

3 Reading the black box: the Interpretable-CHRL program

Before replacing an architecture on interpretability grounds, one owes it the strongest post-hoc reading one can produce. We therefore ran a dedicated program (Interpretable-CHRL) over the

frozen trained policy — no reward change, no re-optimization, no re-fit — asking one question with increasing precision: *which latent structures are actionable under controlled intervention?*

3.1 The staged pipeline

The program proceeds in stages over a frozen rollout (Figure 2, left): log the hidden states and action chunks; *discover* primitives by unsupervised clustering; label them behaviorally from the trading logs; attach outcome diagnostics; and only then intervene. The discovery stage yields a compact vocabulary with healthy statistics: a $K=8$ K-means codebook with utilization 1.000 (every primitive is used), perplexity 6.91 (usage is spread, not collapsed), median run length 18 trading days (primitives are *episodes*, not jitter), and cross-fold normalized mutual information 0.634 (the vocabulary largely replicates on held-out folds). The interventions themselves are one edit deep: take the policy’s own hidden state h_t , push it along an interpretable direction d with strength α , and pass the edited state through the *original* frozen action head,

$$\tilde{h}_t = h_t + \alpha d, \quad \tilde{a}_t \sim \pi_\theta(\cdot \mid \tilde{h}_t), \quad (2)$$

so any behavioral change is attributable to the edit alone. Five intervention families are tested, each phrased as a falsifiable question:

type	question
suppression	does damping a locally harmful primitive improve the counterfactual rollout?
promotion	does amplifying a locally beneficial primitive improve it?
best-context map	does the best primitive depend on market regime (train-only target)?
replacement	does substituting a context-appropriate primitive for a weak one improve it?
repair	do a few targeted, context-conditional edits beat matched random alternatives?

3.2 What post-hoc buys — under controls

The central methodological lesson arrived immediately: on this substrate, a *random-target joint* edit (hidden state + hidden action) also “improves” the frozen 2022–23 rollout (+0.552 pp) — gentle noise perturbs a locally suboptimal policy off its habits. Every claim below is therefore *control-adjusted*: the reported value is the targeted edit’s delta minus its matched-random twin’s (Figure 2, right). Under that discipline, three results survive. First, hidden states are control surfaces: promoting a specific primitive yields +0.540 pp raw but +0.263 pp control-adjusted (+0.049 Sharpe); replacing a weak primitive with its context-best alternative yields +0.421 pp raw, +0.302 pp adjusted. Second, hidden *actions* are too: a learned adapter on the action pathway gives +0.377 pp raw, +0.265 pp adjusted. Third, the effects are specific: adding targeted repairs of two diagnosed primitives to the joint promotion baseline reaches +0.508 pp where the matched random-repair control reaches only +0.316 pp (+0.192 pp of specificity). Post-hoc analysis is not nothing; we state this against our own thesis: it found real, modular, editable structure inside a policy that was never built to have any.

3.3 Where post-hoc stops

The same program measures its own ceiling, and every question we most cared about hit the same wall. *Memory*: the policy’s state is smeared through the network, so no bound exists on what a hidden-state edit can touch — the matched-random result above is exactly the symptom (noise moves behavior about as much as intention; only the *difference* is knowledge). *Commands*: the control surface is ~ 214 engineering knobs, none of which is a semantic command; steering




Figure 2: The Interpretable-CHRL program. Left: the staged pipeline over the frozen policy — log, discover ($K=8$ codebook; utilization 1.000; perplexity 6.91; median run 18d; cross-fold NMI 0.634), label, diagnose, intervene (Eq. 2), control. Right: the flagship counterfactuals on the frozen 2022–23 rollout, raw vs control-adjusted; the dotted line is the random-target joint edit (+0.552 pp) that makes control adjustment mandatory.

means forcing the latent toward a post-hoc K-means centroid and checking that behavior moves *consistently* — a correlation, not an intervention. *Story*: when we scored the deployed agent’s own stance log on the simulatability metric of §6, a $K \leq 9$ named story reproduced only 0.244 of its behavior — the latent is genuinely incompressible at a human-story budget. And *scope*: all counterfactual evidence lives on one frozen, stress-heavy window (2022–23), because the frozen-policy discipline that makes the measurements clean also chains them to a fixed rollout. Transparent layers, opaque core. The program’s own stated next step — “turn the primitive-level signal into a cleaner training-time objective or policy-improvement loop” — is the bridge this paper crosses: instead of discovering primitives after training, *build the agent out of named parts*.

4 CRYSTAL-1: legible by construction

CRYSTAL-1 rebuilds the core around three commitments, each a direct answer to a measured failure of §3.3 (Figure 3).

1. Memory is a named posterior. The agent’s only recurrent state is a belief over K market regimes — a point on the simplex Δ^{K-1} — in the POMDP sense of a belief state [Kaelbling et al., 1998], over Hamilton-style regimes [Hamilton, 1989]. For the $K=2$ US machinery, with regime densities $f_z(r) = \mathcal{N}(r; \mu_z, \sigma_z^2)$, $z \in \{\text{bull}, \text{bear}\}$, and a sticky transition prior $\pi_{zz'}$, the belief is the explicit Hamilton filter

$$b_t = \mathbb{P}(z_t = \text{bear} \mid \mathcal{F}_t) = \frac{f_{\text{bear}}(r_t) \tilde{b}_t}{f_{\text{bear}}(r_t) \tilde{b}_t + f_{\text{bull}}(r_t) (1 - \tilde{b}_t)}, \quad (3)$$

$$\tilde{b}_t = \pi_{BB} b_{t-1} + \pi_{GB} (1 - b_{t-1}),$$

where $\mathcal{F}_t$ is the information available at t (§5). *Named* is load-bearing: because the coordinate has a meaning, a write to the belief is a sentence (“believe bear at 0.8”), and because b_t is the sole memory channel, the blast radius of any memory intervention is bounded by construction — the property Eq. 2 could test for but never guarantee.

2. The policy head is the story. The decision layer is a small tree over $[b_t, \text{inventory}, \text{time}]$. In the soft form [Frosst and Hinton, 2017] each inner node i routes by a sigmoid gate and each




Figure 3: The journey in one picture. CHRL factors decisions into readable layers but keeps an unnamed 64-d latent core with ~ 214 knobs, explained only after training (§3). CRYSTAL-1 makes the explanation the mechanism: a named belief as the sole memory channel, a readable rule as the policy itself, and five named command levers.

leaf ℓ holds a distribution over actions:

$$p_i(x) = \sigma(\beta(w_i^\top x + c_i)), \quad P^\ell(x) = \prod_{i \in \text{path}(\ell)} p_i(x)^{\mathbf{1}[\ell \not\prec i]} (1 - p_i(x))^{\mathbf{1}[i \searrow \ell]}, \quad (4)$$

$$\pi(a \mid x) = \sum_{\ell} P^\ell(x) Q_a^\ell,$$

with the inverse temperature β annealed so the trained tree is effectively hard. The constraint is that the tree *is* the policy — there is no residual network for behavior to hide in — and it must pay no performance tax (§7). At the limit the readable rule degenerates to a sentence; the certified US rule of §7.1 is literally

$$e_t = \begin{cases} 0.74, & b_t > 0.66 \\ 1.00, & \text{otherwise,} \end{cases} \quad (5)$$

where e_t is the equity exposure.

3. Commands are few, named, and governed. The ~ 214 -knob surface collapses to **five agent-facing levers**:

SET_BELIEF · EXPOSURE_MODE · GROUP_CONCENTRATION_CAP · ABSTAIN_FLOOR · DRAWDOWN_BUDGET

(the concentration cap is *hard*, fixing the predecessor’s soft-cap failure). Everything else is engineering-fenced. A write is the *do*-operator of Pearl [2009] made concrete: SET_BELIEF(b^* , w) realizes

$$b_t \leftarrow (1 - w) b_t + w b^*, \quad w \in [0, 1], \quad (6)$$

and a *dose* nudge of δ nats acts in log-odds, $\text{logit}(b_t) \leftarrow \text{logit}(b_t) + \delta$. Writes must pass a non-inferiority bar on the certified objective before they are allowed to persist (§7).

5 Data and evaluation

All quantitative claims are computed on real market data under point-in-time (PIT) discipline. The universe is daily total-return series for six liquid US-listed assets plus cash — SPY (US large-cap equities), EFA (developed ex-US), IEF and TLT (7–10y and 20y+ Treasuries), TIP (inflation-linked), GLD (gold), with the 13-week T-bill rate as the cash leg — split- and dividend-adjusted

closes converted to total-return relatives, the cash leg accruing $r_t^f = y_t/252$. The exposure-rule experiments (§7) additionally use the lab’s Dow-30 daily panel (2000–2026). Every panel enters through a signed data gate (file hash, source URL, vintage recorded before any experiment may read it) and is ground-truthed against an independent source — a discipline bought by experience: a Chinese panel once passed every internal consistency check while carrying a mixed clock (close and volume lagged a day against high and low); only external ground-truthing caught it, and every result computed on that panel was voided. A decision issued at t may use only $\mathcal{F}_t = \sigma(\{x_s : s \leq t\})$, with publication lags respected; evaluation stands at historical origins and is scored on realized continuations, never re-fit on them.

Two evaluation formulas recur. For a daily excess-return sequence x_t and wealth path W_t :

$$\text{Sharpe} = \frac{\sqrt{252} \bar{x}}{\hat{\sigma}(x)}, \quad \text{max drawdown} = \min_t \left(\frac{W_t}{\max_{s \leq t} W_s} - 1 \right). \quad (7)$$

The goal-planning machinery that CRYSTAL-1 serves in production — scenario kernels by stationary bootstrap [Politis and Romano, 1994], dynamic programming on the funding ratio [Das et al., 2020, Browne, 1999], and the promise backtests — is specified and evaluated in the companion paper (*Hello Crystal*); here we use only its policy *tables*, which are the objects the simulatability metric reads.

6 CrystalScore: measuring “legible” so it can lose

A legibility claim that cannot lose is marketing. We score every architecture with

$$\text{CRYSTALSCORE} = F \times S \times St, \quad (8)$$

at a fixed parsimony budget $K \leq 9$ (at most nine named rules/leaves), with each axis operationalized so that it *could have* come out low [Doshi-Velez and Kim, 2017, Jacovi and Goldberg, 2020, Lipton, 2018]:

$$F \text{ (faithfulness)} = \mathbb{E} \left[\mathbf{1} [\pi(b \leftarrow \text{do}(b^*)) = \pi^*(b^*)] \right], \quad (9)$$

$$S \text{ (simulatability)} = \text{acc}(\text{story}_K, \pi) \quad \text{s.t.} \quad |\mathcal{R}(\text{story}_K) - \mathcal{R}(\pi)| \leq \varepsilon, \quad (10)$$

$$St \text{ (stability)} = \frac{1}{n} \sum_{i=1}^n \mathbf{1} [\text{mechanism replicates on seed } i]. \quad (11)$$

Faithfulness (Eq. 9): *do*-writes must land, judged against a reference optimum, not against the policy’s own habits. Simulatability (Eq. 10): a K -rule story must match the policy *at performance parity* $\mathcal{R}$ — a story that reproduces a lobotomized policy is cheating. Stability (Eq. 11): the mechanism must replicate across independent training seeds. The non-saturating companion is the **MDL deficit** [Rissanen, 1978],

$$\Delta_{\text{MDL}} = \text{Acc}(\pi) - \text{Acc}(\text{story}_K), \quad (12)$$

the accuracy lost by forcing the explanation to be short; 0 means a short story loses nothing.

A measurement correction, reported as a finding. Faithfulness probes can silently lie. Our original probe applied fixed-size belief writes (+0.2/ +0.4) and checked whether the action followed. On threshold policies whose belief input saturates, such writes never cross the decision boundary — so a policy that *is* the rule by construction (a behavior-cloned distillation with accuracy 0.97) read $F=0.00$. The repaired protocol uses matched-state *contrast writes* ($b=0.2$ vs $b=0.8$, spanning both certified thresholds). Under the corrected probe, cold-trained RL heads




Figure 4: The identical scalar on both architectures, real data only: the CHRL line is scored on its own deployed-panel stance log; CRYSTAL-1 on the US goal machinery (tree-vs-champion simulatability) and the certified Dow rule (10-seed stability, write fidelity). Faithfulness is 1.0 for both — steering works on a black box too; the discriminating axis is whether a $K \leq 9$ named story can reproduce behavior at parity.

remain unfaithful under both probes (F strict 0.00, contrast 0.00), while teacher-warm-started heads score contrast-write $F = 0.77\text{--}0.83$ where the strict probe read 0.03 (§7). The lesson generalizes: *an interpretability metric is an instrument, and instruments need calibration surfaces* — this one was caught on a designed environment where ground truth was known by construction (§8).

Real-data values. On its own deployed-panel stance log, the CHRL-line agent scores $F=1.00$ (latent steering does move behavior — faithfulness alone cannot distinguish the designs), $S=0.244$ (§3.3), $St=0.619$: CRYSTALSCORE=0.151. CRYSTAL-1’s axes on the real US machinery: $S=0.92$ (an 8-leaf tree reproduces the goal-planner champion’s policy table at parity); $St=1.00$ (the certified rule’s behavior replicates on 10/10 training seeds); $F=1.00$ (writing the bear belief past the threshold changes exposure exactly as named); $\Delta_{MDL}=0.0$: CRYSTALSCORE=0.92 (Figure 4). The entire $6\times$ gap is Simulatability — exactly what “born legible” buys. Alpha does not enter the score.

7 The comparison: influence, certification, and the final table

A legible agent is only useful if you can *change* it on purpose. We exercised four families of influence on CRYSTAL-1’s PPO reward-dial heads — a research prototype, not the champion — on the real US daily panel (train 2010–18, held-out evaluation after), each with a control battery so we report what the levers actually do, not what they appear to (Figure 5).

R — reward shaping. The risk-targeted reward

$$r_t = e_t x_t - \lambda_R \max(0, DD_t - D^*) \quad (13)$$

(e_t exposure, x_t the asset return, DD_t running drawdown, D^* the budget) moves the exposure dial, but the penalty *alone does not enforce the budget*: only the 5% head holds its budget (−4.81% held-out max drawdown, by parking in cash) while the 8% and 12% heads breach theirs (−15.9%). A drawdown-matched twin shows $z = 0.00$ timing skill: reward shaping buys




Figure 5: The four influence levers on the real daily panel, under controls. R: the reward penalty holds only the 5% budget — the penalty does not enforce the budget. T: the teacher warm-start cures the cold near-uniform collapse ($0.21 \rightarrow 0.73$; corrected contrast-write $F = 0.77\text{--}0.83$). I: a matched-random leaf moves the cold head as much as the targeted one — only the non-inferiority *refusal* (Eq. 14) survives. A: the cold head is a legible, stable dial ($S=1.0$, $St=1.0$, $F=0.03$), not a faithful command surface.

risk placement, not timing. Enforcing the budget requires a separate constrained objective (a Lagrangian head drives the breach down at a 15–37 pp goal-probability cost) — one more reason the transparent planner, which respects capacity by a hard pre-filter, stays champion.

T — teachers. Trained cold, the heads’ argmax behavior looked sensible while the underlying preferences were nearly flat — an effectively untrained network masked by deterministic evaluation (mean max action-probability $0.21 \approx 1/5$). The cure is behavior cloning from the certified rule (Eq. 5), $L_{BC}(\theta) = -\mathbb{E}_{(s,a^*)}[\log \pi_\theta(a^* | s)]$, 87% agreement, used as a warm start [Ross et al., 2011, Rusu et al., 2016]: mean max action-probability $0.21 \rightarrow 0.73$, the belief becomes load-bearing, and the corrected contrast-write probe (§6) shows the teacher installs the rule’s *causal* belief→action structure ($F = 0.77\text{--}0.83$), which PPO fine-tuning does not erode on this substrate. Cold discovery still fails; teaching works. (Leakage note: this teacher was *selected* on the 2022–23 hold, so its hold read is contaminated; the leak-free teacher in the pipeline is the planner champion, and a fresh read is pre-registered for 2027.)

I — interventions, under controls. The honest result overturns a first-pass claim. On the cold reward-dial heads, belief writes and per-primitive promotions *appear* to steer behavior — but a matched-random leaf moves behavior exactly as much as the targeted one (the placebo fails), and belief-write fidelity is 3%: the apparent steering was a near-uniform-head artifact, precisely the failure mode §3.2’s controls exist to catch. What survives is the refusal machinery: the certification gate correctly refuses a +1-nat promotion that fails the non-inferiority test

$$\text{accept the write} \iff \hat{J}_{\text{post}} \geq \hat{J}_{\text{pre}} - \delta_{\text{NI}}, \quad (14)$$

and likewise rejects on the hold a reward “winner” selected on development ($z = 1.34 < 2.15$).

A — architecture. Tree depth is a legibility dial set before training: depth-2 yields a 4-leaf policy, depth-3 the 8-leaf one. Scoring the prototype head itself: perfectly simulatable ($S=1.0$; a 4-leaf story distills the 8-leaf head exactly) and stable ($St=1.0$, 3 seeds) but — when cold — not faithful ($F=0.03$; CRYSTALSCORE=0.03): a legible, stable *dial*, not a command surface. The faithful command surface is either taught (lever T) or built (the champion’s rule, $F=1.0$).

7.1 A certified readable rule on untouched data

The self-improvement loop’s frozen certification gate (specified in the companion paper) has certified exactly one real-data policy to date — and it is a sentence (Eq. 5). It entered with 7 accepts and 0 placebo accepts, and held one-shot on 629 untouched trading days: Sharpe 1.46 vs 1.39 buy-and-hold, max drawdown -12.9% vs -15.5% , risk-adjusted improvement $z = +2.06$. The claim is a *risk* claim (same return, meaningfully less pain), certified as such; a written return-improvement rule was separately tested and closed (-8.8%). The rule’s machinery was then audited on the real panel: naming faithfulness 1.0 (the state called “bear” is the state in which de-risking is optimal under the world model); 10-seed behavioral stability 1.0; rule re-discovery from scratch 3/3; belief MDL deficit 0.0. These are the CRYSTAL-1 entries of Figure 4 and Table 1.

	CHRL	CRYSTAL-1
Memory	64-d unnamed latent, smeared; blast radius unboundable	named posterior b_t (Eq. 3); the sole memory channel
Policy legibility	post-hoc codes over a continuous latent; story acc. 0.244	the policy <i>is</i> a table/rule; 8-leaf story acc. 0.92 at parity
Steering primitive	force the latent toward a K-means centroid (correlation; matched-random controls mandatory)	named <i>do-write</i> (Eq. 6), fidelity 1.0
Refusals	OOD gate shrinks or zeroes steps	the NI bar (Eq. 14) refused a +1-nat promotion
Control surface	~ 214 knobs	5 named levers
Seed stability	0.619	1.00 (10 seeds)
CRYSTALSCORE (Eq. 8)	0.151	0.92
Flagship real-data result	walk-forward: Sharpe 1.09, maxDD -7.34% ; control-adjusted targeted edits $+0.19$ – 0.30 pp	certified rule, 629 untouched days: Sharpe 1.46 vs 1.39, maxDD -12.9% vs -15.5%
Deployment status	deployed: live firewall, operating history	goal machinery live on the US universe (paper track); client certification still <i>refused</i> , with dated unlocks

Table 1: The final comparison, real data only. The honest split: CHRL remains the longer-operating deployed system, and its post-hoc program found real editable structure; CRYSTAL-1 wins every legibility and control axis and carries its own real-data results, while its client-facing certification correctly refuses until the pre-registered reads mature.

8 The designed-market experiment: can a market reward only a clean mind?

Real markets have very noisy logic: the regime is largely priced, so the honest value of even a perfect regime belief can be zero — and a null result on real data cannot distinguish “the instrument is broken” from “there was nothing to detect.” This motivates one deliberately synthetic experiment: *can one build a market that rewards only a “clean” mind — an environment where the belief information carries value by construction, so that command-surface machinery can be calibrated at all?*

We built such an environment (a regime market whose returns are generated from latent states the belief can, in principle, track) and used it strictly as **instrument calibration**. On it: belief writes were *proved* causal against an exogenous belief-MDP optimum (agreement 0.80/0.835 for the two heads); a likelihood-based *lie detector* separates sharp-but-wrong commands from




Figure 6: The designed-market experiment and its fence. Left: the environment built to reward a clean mind — belief information worth +282% of blind by construction, so command machinery (causal writes, the lie detector, composition governance, probe calibration) can be calibrated. Right: the same measurement on a real book returns zero, and the gate closes.

uncertain-but-right ones (excess predictive NLL 0.43 vs ≈ 0 ; detection 0.71 at 5% false alarms) while the naïve entropy gate points *backwards* (it trusts the confident lie); a 216-cell factorial showed command composition is non-additive with *sign-flipping* cross-terms (median $|\varepsilon| = 2.66$) — which is why the write surface is governed by a calibrated ledger, a co-activation cap, a kNN manifold governor (24/24 engineered off-manifold corners flagged at 1.2% held-out FPR) and a writ ladder with a first false-accept estimate (0/48; $\lesssim 6\%$ by the rule of three). On a *hard* designed substrate the gate also certified a genuine return $\leftrightarrow$ legibility *tension* — legibility is not free everywhere, and anyone claiming otherwise is measuring easy tasks. It was also this environment that exposed the faithfulness-probe blindness of §6.

The bridge back to reality is an explicit **value-of-information gate** (Figure 6): the designed market shows $\text{Vol} = 5.92$ (+282% of a belief-blind baseline) and the gate *opens*; on a real book the measured Vol is 0 and the gate *closes* — so nothing calibrated on the designed market is allowed to generate a real-market claim, and no designed-market number appears in §7.

9 What did not work, and what bounds the claims

A paper about honest machinery should list its own falsifications; each cost us a believed hypothesis, and the machinery that killed them is part of the contribution. (i) **Raw counterfactual lifts were inflated $\sim 2\times$** : a random-target joint edit “improves” the frozen rollout by +0.552 pp (§3.2) — without matched controls, every post-hoc intervention result in §3 would have been overstated. (ii) **The strict faithfulness probe read $F=0.00$ on a policy that was the rule by construction** (§6) — an uncalibrated interpretability metric lied in the pessimistic direction. (iii) **Cold RL heads are not command surfaces**: the placebo fails, write fidelity is 3%; only teaching or building installs causal structure (§7). (iv) **A written return-improvement rule closed (-8.8%)**: the certified claim is risk-shaped, not alpha-shaped. (v) **A PPO challenger loses to the readable planner table** on the same action space (regret 6.6 pp) — reported because the field’s default is the opposite claim. (vi) **Legibility is provably not free** on a hard designed substrate (the certified tension of §8). (vii) **Post-hoc interpretability was not worthless**: it found the load-bearing concepts (motifs, codes, counterfactual replay) that CRYSTAL-1 promoted from probes into architecture — the right summary is that building them in made the story $\sim 4\times$ more compressible (0.244 $\rightarrow$ 0.92 at the same budget).

10 Related work

Interpretability-by-construction vs post-hoc explanation [Rudin, 2019, Lipton, 2018]; rigor and evaluation of interpretability [Doshi-Velez and Kim, 2017, Hoffman et al., 2018]; faithfulness as a criterion [Jacovi and Goldberg, 2020]. CRYSTAL-1’s named belief is a concept-bottleneck design in the sense of Koh et al. [2020] — the belief is a supervised, human-named intermediate through which all memory must pass — with the difference that our bottleneck is also the *write* surface, so faithfulness is tested by intervention (Eq. 9), not only by probe accuracy; and where circuit-style analysis [Olah et al., 2020] reverse-engineers named structure *out of* trained networks, CrystalRL builds the named structure *in* and reserves reverse-engineering for the audit. Belief-state control [Kaelbling et al., 1998] and regime models [Hamilton, 1989] supply the named memory; soft trees [Frosst and Hinton, 2017] the legible head; MDL [Rissanen, 1978] the non-saturating ruler; interventional causality [Pearl, 2009] the faithfulness protocol. Discretized latent-action behavior generation [Lee et al., 2024] pursues the complementary direction — learned discrete codes for *expressiveness* — where CrystalRL fixes the codes to *named* coordinates for legibility. Hierarchical RL as partial legibility [Sutton et al., 1999, Bacon et al., 2017]; PPO and constrained variants [Schulman et al., 2017, Achiam et al., 2017]; goal-conditioned value functions [Schaul et al., 2015]; imitation-then-optimize [Ross et al., 2011, Rusu et al., 2016]; reward shaping [Ng et al., 1999]; goal-based planning [Das et al., 2020, Browne, 1999]; the stationary bootstrap [Politis and Romano, 1994]. The self-improvement layer — LLM proposers over verified evaluators [Romera-Paredes et al., 2024, Novikov et al., 2025, Ma et al., 2024, Wang et al., 2023] and the statistics of adaptive search over noisy financial evaluators [Bailey and López de Prado, 2014] — is treated in the companion paper (*Hello Crystal*); this paper builds the object those loops are allowed to modify.

11 Conclusions: what the journey establishes about interpretable RL

Five findings travel beyond this system.

1. The discriminating axis is simulatability, not faithfulness. Steering “works” on a black box too ($F=1.0$ for both architectures); what born-legible buys is that a $K \leq 9$ named story reproduces behavior at performance parity ($0.244 \rightarrow 0.92$). Evaluations of interpretable RL that report only intervention success are measuring the wrong axis.

2. Post-hoc analysis finds structure but cannot bound it. The strongest post-hoc program we could run found a real, complete, editable primitive vocabulary — and still could not bound the blast radius of an edit, name a command, or compress the policy into a story. The two families are not rivals: post-hoc discovery told us *which* concepts to build into the successor.

3. Controls and calibration are not optional. A random-target edit improves the rollout (+0.552 pp); a placebo leaf moves a cold head as much as the targeted one; a strict faithfulness probe reads 0.00 on a policy that is the rule by construction. Every interpretability measurement in this paper that lacked a control or a calibration surface was wrong on the first pass — always in a direction favorable to a good story.

4. Causal command structure is taught or built, never free. Cold RL training produced legible, stable *dials* with near-zero faithfulness; behavior cloning from a certified readable rule installed the causal belief→action structure ($F = 0.77\text{--}0.83$) and PPO fine-tuning did not erode it. If you need a command surface, put it there on purpose.

5. Legibility is not free everywhere — and where it is expensive is informative. On an easy substrate the readable rule pays no tax; on a hard designed substrate a certified return↔legibility tension appears. Mapping *where* the tension binds is a better research question than asserting either side of it.

The forward-looking program is falsifiable by design; its flagship hypotheses: machine simulatability predicts *human* simulatability (kill: humans at chance while the machine metric is at ceiling); a legibility scaling law — optimal vocabulary size K tracks substrate behavioral complexity (kill: no relation across universes); interpretability *drift* — naming faithfulness decays under a self-improvement loop unless re-certified (kill: F stays 1.0 over 20+ accepted changes); motif conservation — CHRL’s discovered primitives re-emerge as CRYSTAL-1 rule clauses (kill: no better-than-chance correspondence); dose-response of SET_BELIEF is monotone and saturating on certified substrates (kill: non-monotone regions); the composition ledger’s cross-terms appear on the real panel’s visited manifold (kill of the premise: on-manifold interactions are additive); and CrystalScore transfers to the *improver* — the coding agent’s stated rationales predict its measured effects (kill: rationale–effect agreement at chance, itself a publishable confabulation result).

12 Limitations

(i) The scenario-based probabilities that the goal machinery quotes remain uncalibrated research estimates; the calibration gates and their dated forward reads (2027-07-12; an untouched multi-year fold maturing 2029) live in the companion paper. (ii) The real-data CrystalScore’s simulatability axis is measured on the planner champion and the certified rule family — not yet on every head. (iii) The command-surface *proofs* (causal writes against an exogenous optimum, the lie detector, composition governance) are designed-market calibrations (§8); their real-data analogues are open hypotheses. (iv) All Interpretable-CHRL counterfactuals live on one frozen, stress-heavy window (2022–23) by design. (v) The certified rule is one rule family on one substrate; hardening across substrates, seeds and profiles is the program’s first milestone. (vi) Execution economics beyond 10 bp switching costs (taxes, spreads, gaps) are only partially bounded.

Acknowledgments and drafting discipline

Prepared *write-paper-first* [Eisner, 2010]. Experiments were run and logged under the repository’s honesty contract — every entry names its null, its seed, and its worst caveat — in the spirit of Feynman [1974]: “*The first principle is that you must not fool yourself — and you are the easiest person to fool.*” LLM coding agents (Claude, Codex) operated the experimental loops under the stated controls; the agents propose, the gates dispose. The author thanks Joseph Lynch for discussions on the interpretability track.

References

- Joshua Achiam, David Held, Aviv Tamar, and Pieter Abbeel. Constrained policy optimization. In *ICML*, 2017.
- Pierre-Luc Bacon, Jean Harb, and Doina Precup. The option-critic architecture. In *AAAI*, 2017.
- David H. Bailey and Marcos López de Prado. The deflated Sharpe ratio: correcting for selection bias, backtest overfitting, and non-normality. *Journal of Portfolio Management*, 40:94–107, 2014.
- Sid Browne. Reaching goals by a deadline: digital options and continuous-time active portfolio management. *Advances in Applied Probability*, 31:551–577, 1999.
- Sanjiv R. Das, Daniel Ostrov, Anand Radhakrishnan, and Deep Srivastav. Dynamic portfolio allocation in goals-based wealth management. *Computational Management Science*, 17:613–640, 2020.

- Finale Doshi-Velez and Been Kim. Towards a rigorous science of interpretable machine learning. *arXiv:1702.08608*, 2017.
- Jason Eisner. Write the paper first. <https://www.cs.jhu.edu/~jason/advice/write-the-paper-first.html>, 2010.
- Richard P. Feynman. Cargo cult science. Caltech commencement address, 1974.
- Nicholas Frosst and Geoffrey Hinton. Distilling a neural network into a soft decision tree. *arXiv:1711.09784*, 2017.
- James D. Hamilton. A new approach to the economic analysis of nonstationary time series and the business cycle. *Econometrica*, 57:357–384, 1989.
- Robert R. Hoffman, Shane T. Mueller, Gary Klein, and Jordan Litman. Metrics for explainable AI: challenges and prospects. In *arXiv:1812.04608*, 2018.
- Alon Jacovi and Yoav Goldberg. Towards faithfully interpretable NLP systems: how should we define and evaluate faithfulness? In *ACL*, 2020.
- Leslie Pack Kaelbling, Michael L. Littman, and Anthony R. Cassandra. Planning and acting in partially observable stochastic domains. *Artificial Intelligence*, 101:99–134, 1998.
- Pang Wei Koh, Thao Nguyen, Yew Siang Tang, Stephen Mussmann, Emma Pierson, Been Kim, and Percy Liang. Concept bottleneck models. In *ICML*, 2020.
- Seungjae Lee, Yibin Wang, Haritheja Etukuru, H. Jin Kim, Nur Muhammad Mahi Shafiullah, and Lerrel Pinto. Behavior generation with latent actions. In *ICML*, 2024. arXiv:2403.03181.
- Zachary C. Lipton. The mythos of model interpretability. *Communications of the ACM*, 61:36–43, 2018.
- Yecheng Jason Ma et al. Eureka: human-level reward design via coding large language models. In *ICLR*, 2024.
- Volodymyr Mnih, Koray Kavukcuoglu, David Silver, et al. Human-level control through deep reinforcement learning. *Nature*, 518:529–533, 2015.
- Andrew Y. Ng, Daishi Harada, and Stuart Russell. Policy invariance under reward transformations: theory and application to reward shaping. In *ICML*, 1999.
- Alexander Novikov et al. AlphaEvolve: a coding agent for scientific and algorithmic discovery. *DeepMind technical report*, 2025.
- Chris Olah, Nick Cammarata, Ludwig Schubert, Gabriel Goh, Michael Petrov, and Shan Carter. Zoom in: an introduction to circuits. *Distill*, 2020. doi:10.23915/distill.00024.001.
- Judea Pearl. *Causality: Models, Reasoning, and Inference*. Cambridge University Press, 2nd edition, 2009.
- Dimitris N. Politis and Joseph P. Romano. The stationary bootstrap. *Journal of the American Statistical Association*, 89:1303–1313, 1994.
- Jorma Rissanen. Modeling by shortest data description. *Automatica*, 14:465–471, 1978.
- Bernardino Romera-Paredes et al. Mathematical discoveries from program search with large language models. *Nature*, 625:468–475, 2024.

- Stéphane Ross, Geoffrey Gordon, and Drew Bagnell. A reduction of imitation learning and structured prediction to no-regret online learning. In *AISTATS*, 2011.
- Cynthia Rudin. Stop explaining black box machine learning models for high stakes decisions and use interpretable models instead. *Nature Machine Intelligence*, 1:206–215, 2019.
- Andrei A. Rusu et al. Policy distillation. In *ICLR*, 2016.
- Tom Schaul, Daniel Horgan, Karol Gregor, and David Silver. Universal value function approximators. In *ICML*, 2015.
- John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms. *arXiv:1707.06347*, 2017.
- Richard S. Sutton, Doina Precup, and Satinder Singh. Between MDPs and semi-MDPs: a framework for temporal abstraction in reinforcement learning. *Artificial Intelligence*, 112: 181–211, 1999.
- Guanzhi Wang et al. Voyager: an open-ended embodied agent with large language models. *arXiv:2305.16291*, 2023.